

Лабораторная работа 4
по дисциплине
«Методы машинного обучения»
на тему
«Подготовка обучающей и тестовой выборки,
кросс-валидация и подбор гиперпараметров на
примере метода ближайших соседей.»

Выполнил:
студент группы ИУ5-21М
Могильников И. А.

1. Описание задания

Цель лабораторной работы: изучение сложных способов подготовки выборки и подбора гиперпараметров на примере метода ближайших соседей.

2. Задание

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите модель ближайших соседей для произвольно заданного гиперпараметра K . Оцените качество модели с помощью трех подходящих для задачи метрик.
5. Постройте модель и оцените качество модели с использованием кросс-валидации. Проведите эксперименты с тремя различными стратегиями кросс-валидации.
6. Произведите подбор гиперпараметра K с использованием `GridSearchCV` и кросс-валидации.
7. Повторите пункт 4 для найденного оптимального значения гиперпараметра K . Сравните качество полученной модели с качеством модели, полученной в пункте 4.
8. Постройте кривые обучения и валидации.

3. Ход выполнения лабораторной работы

3.1. Выбор датасета

В качестве исходных данных выбираем датасет Heart Disease UCI (<https://www.kaggle.com/ronitf/heart-disease-uci>). 303 записи, 14 признаков, целевой признак относится к наличию болезни сердца у пациента: 0 - нет болезни сердца, 1 - есть.

```
In [0]: from google.colab import drive, files
        drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`

```
In [0]: from google.colab import files
import os
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
os.listdir()
data = pd.read_csv('drive/My Drive/mmo_datasets/heart.csv',
                  sep=";", encoding="iso-8859-1")
```

```
In [0]: total_count = data.shape[0]
num_cols = []
for col in data.columns:
    # Количество пустых значений
    temp_null_count = data[data[col].isnull()].shape[0]
    dt = str(data[col].dtype)
    if temp_null_count>0:
        num_cols.append(col)
        temp_perc = round((temp_null_count / total_count) * 100.0, 2)
        print('Колонка {}. Тип данных {}. Количество пустых значений {}, {}%.'
              .format(col, dt, temp_null_count, temp_perc))
```

```
data_cleared = data
```

```
In [0]: uniquevalues = np.unique(data_cleared['target'].values)
uniquevalues
```

```
Out[0]: array([0, 1])
```

```
In [0]: data_cleared.head(10)
```

```
Out[0]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope \
0	63	1	3	145	233	1	0	150	0	2.3	0
1	37	1	2	130	250	0	1	187	0	3.5	0
2	41	0	1	130	204	0	0	172	0	1.4	2
3	56	1	1	120	236	0	1	178	0	0.8	2
4	57	0	0	120	354	0	1	163	1	0.6	2
5	57	1	0	140	192	0	1	148	0	0.4	1
6	56	0	1	140	294	0	0	153	0	1.3	1
7	44	1	1	120	263	0	1	173	0	0.0	2
8	52	1	2	172	199	1	1	162	0	0.5	2
9	57	1	2	150	168	0	1	174	0	1.6	2

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1
5	0	1	1
6	0	2	1
7	0	3	1
8	0	3	1
9	0	2	1

3.2. train_test_split

```
In [0]: target = data_cleared['target']
data_cleared = data_cleared.drop('target', axis=1)
```

```
In [0]: data_cleared.head(10)
```

```
Out[0]: age sex cp trestbps chol fbs restecg thalach exang oldpeak slope \
0 63 1 3 145 233 1 0 150 0 2.3 0
1 37 1 2 130 250 0 1 187 0 3.5 0
2 41 0 1 130 204 0 0 172 0 1.4 2
3 56 1 1 120 236 0 1 178 0 0.8 2
4 57 0 0 120 354 0 1 163 1 0.6 2
5 57 1 0 140 192 0 1 148 0 0.4 1
6 56 0 1 140 294 0 0 153 0 1.3 1
7 44 1 1 120 263 0 1 173 0 0.0 2
8 52 1 2 172 199 1 1 162 0 0.5 2
9 57 1 2 150 168 0 1 174 0 1.6 2
```

```
ca thal
0 0 1
1 0 2
2 0 2
3 0 2
4 0 2
5 0 1
6 0 2
7 0 3
8 0 3
9 0 2
```

```
In [0]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(
    data_cleared,
    target,
    test_size=0.2,
    random_state=1
)
```

```
In [0]: X_train.shape, Y_train.shape
```

```
Out[0]: ((242, 13), (242,))
```

```
In [0]: X_test.shape, Y_test.shape
```

```
Out[0]: ((61, 13), (61,))
```

3.3. Обучение для произвольного параметра K

```
In [0]: from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
```

```
In [0]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, Y_train)
predicted = knn_model.predict(X_test)
predicted
```

```
Out[0]: array([0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0,
0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0,
1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0])
```

```
In [0]: from sklearn.metrics import accuracy_score
        accuracy_score(Y_test, predicted)
```

```
Out[0]: 0.5737704918032787
```

```
In [0]: from sklearn.metrics import balanced_accuracy_score
        balanced_accuracy_score(Y_test, predicted)
```

```
Out[0]: 0.5720430107526882
```

```
In [0]: # https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix
        from sklearn.utils.multiclass import unique_labels
        def plot_confusion_matrix(y_true, y_pred, classes,
                                   normalize=False,
                                   title=None,
                                   cmap=plt.cm.Blues):
            """
            This function prints and plots the confusion matrix.
            Normalization can be applied by setting `normalize=True`.
            """
            if not title:
                if normalize:
                    title = 'Normalized confusion matrix'
                else:
                    title = 'Confusion matrix, without normalization'

            # Compute confusion matrix
            cm = confusion_matrix(y_true, y_pred)
            # Only use the labels that appear in the data
            classes = classes[unique_labels(y_true, y_pred)]
            if normalize:
                cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
                print("Normalized confusion matrix")
            else:
                print('Confusion matrix, without normalization')

            fig, ax = plt.subplots()
            im = ax.imshow(cm, interpolation='nearest', cmap=cmap)
            ax.figure.colorbar(im, ax=ax)
            # We want to show all ticks...
            ax.set(xticks=np.arange(cm.shape[1]),
                   yticks=np.arange(cm.shape[0]),
                   # ... and label them with the respective list entries
                   xticklabels=classes, yticklabels=classes,
                   title=title,
                   ylabel='True label',
                   xlabel='Predicted label')

            # Rotate the tick labels and set their alignment.
            plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
                     rotation_mode="anchor")
```

```

# Loop over data dimensions and create text annotations.
fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, format(cm[i, j], fmt),
                ha="center", va="center",
                color="white" if cm[i, j] > thresh else "black")
fig.tight_layout()
return ax

```


```

In [0]: plot_confusion_matrix(Y_test, predicted,
                             classes=np.array(['E', 'E10+', 'M', 'T']),
                             title='Confusion matrix, without normalization')

```

Confusion matrix, without normalization

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f13d7451240>



```

In [0]: from sklearn.metrics import precision_score, recall_score, f1_score
        (precision_score(Y_test, predicted, average='weighted'),
         recall_score(Y_test, predicted, average='weighted'))

```

Out[0]: (0.5753212228622065, 0.5737704918032787)

```

In [0]: f1_score(Y_test, predicted, average='weighted')

```

Out[0]: 0.5688953176899175

3.4. Построение модели и оценка с помощью кросс-валидации

```
In [0]: from sklearn.model_selection import KFold, ShuffleSplit, StratifiedShuffleSplit
        from sklearn.model_selection import cross_val_score, cross_validate
```

```
scoring = {'precision': 'precision_weighted',
           'recall': 'recall_weighted',
           'f1': 'f1_weighted'}
```

```
In [0]: scores1 = cross_validate(KNeighborsClassifier(n_neighbors=2),
                                data_cleared,
                                target,
                                scoring=scoring,
                                cv=KFold(n_splits=3),
                                return_train_score=True
                                )

scores1
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/classification.py:1145: UndefinedMetricWarning:
'recall', 'true', average, warn_for)
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/classification.py:1145: UndefinedMetricWarning:
'recall', 'true', average, warn_for)
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/classification.py:1145: UndefinedMetricWarning:
'recall', 'true', average, warn_for)
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/classification.py:1145: UndefinedMetricWarning:
'recall', 'true', average, warn_for)
```

```
Out[0]: {'fit_time': array([0.0035882 , 0.00298381, 0.00277495]),
         'score_time': array([0.01822829, 0.01480269, 0.0153482 ]),
         'test_f1': array([0.28813559, 0.55414336, 0.62585034]),
         'test_precision': array([1.         , 0.69316227, 1.         ]),
         'test_recall': array([0.16831683, 0.56435644, 0.45544554]),
         'train_f1': array([0.80387838, 0.79945524, 0.89967881]),
         'train_precision': array([0.86178676, 0.86071429, 0.93169995]),
         'train_recall': array([0.82673267, 0.80693069, 0.89108911])}
```

```
In [0]: scores2 = cross_validate(KNeighborsClassifier(n_neighbors=2),
                                data_cleared,
                                target,
                                scoring=scoring,
                                cv=ShuffleSplit(n_splits=5, test_size=0.25),
                                return_train_score=True
                                )

scores2
```

```
Out[0]: {'fit_time': array([0.00568986, 0.00283527, 0.00285411, 0.00288081, 0.0028367 ]),
         'score_time': array([0.01616049, 0.01285553, 0.01286197, 0.01299715, 0.0131793 ]),
         'test_f1': array([0.59090453, 0.43465982, 0.54309958, 0.49336384, 0.54641813]),
         'test_precision': array([0.65233425, 0.50489204, 0.55322831, 0.58439201, 0.62388664]),
         'test_recall': array([0.59210526, 0.44736842, 0.55263158, 0.52631579, 0.55263158]),
         'train_f1': array([0.77611602, 0.82921815, 0.78953072, 0.80197881, 0.77611602])}
```

```
'train_precision': array([0.85151099, 0.87616921, 0.85960413, 0.86379331, 0.85151099]),
'train_recall': array([0.78414097, 0.83259912, 0.79295154, 0.8061674 , 0.78414097])}
```

```
In [0]: scores3 = cross_validate(KNeighborsClassifier(n_neighbors=2),
                                data_cleared,
                                target,
                                scoring=scoring,
                                cv=StratifiedShuffleSplit(n_splits=5, test_size=0.2),
                                return_train_score=True
                                )

scores3
```

```
Out[0]: {'fit_time': array([0.00362515, 0.00277042, 0.00277781, 0.00282669, 0.00318956]),
'score_time': array([0.01695871, 0.01117134, 0.01135397, 0.01128387, 0.01146054]),
'test_f1': array([0.61222806, 0.60401357, 0.61928718, 0.5710147 , 0.53801583]),
'test_precision': array([0.66323471, 0.62287796, 0.64371954, 0.58848816, 0.55409836]),
'test_recall': array([0.62295082, 0.60655738, 0.62295082, 0.57377049, 0.54098361]),
'train_f1': array([0.81073333, 0.78385644, 0.79288886, 0.77931087, 0.80184646]),
'train_precision': array([0.86803519, 0.85601355, 0.85991995, 0.85409652, 0.86392588]),
'train_recall': array([0.81404959, 0.7892562 , 0.79752066, 0.78512397, 0.80578512])}
```

```
In [0]: print("%s, %s, %s" % (np.mean(scores1["test_precision"]),
                               np.mean(scores2["test_precision"]),
                               np.mean(scores3["test_precision"])))
```

```
0.8977207579912921, 0.5837466496118905, 0.6144837452406424
```

Лучшую точность модели получилось достичь с использованием стратегии кросс-валидации KFold.

3.5. Подбор гиперпараметра K с использованием GridSearchCV и кросс-валидации

```
In [0]: n_range = np.array(range(2,32,2))
tuned_parameters = [{'n_neighbors': n_range}]
tuned_parameters
```

```
Out[0]: [{'n_neighbors': array([ 2,  4,  6,  8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30])}]
```

```
In [0]: from sklearn.model_selection import GridSearchCV
```

```
clf_gs = GridSearchCV(KNeighborsClassifier(),
                      tuned_parameters,
                      cv=ShuffleSplit(n_splits=5, test_size=0.25),
                      scoring='accuracy')
clf_gs.fit(X_train, Y_train)
```

```
Out[0]: GridSearchCV(cv=ShuffleSplit(n_splits=5, random_state=None, test_size=0.25, train_size=0.75),
                    error_score='raise-deprecating',
                    estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=5, p=2,
```



```

        weights='uniform'),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid=[{'n_neighbors': array([ 2,  4,  6,  8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30])}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='accuracy', verbose=0)

```

In [0]: `clf_gs.best_params_`


Out[0]: `{'n_neighbors': 8}`

In [0]: `clf_gs.best_estimator_`

Out[0]: `KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski', metric_params=None, n_jobs=None, n_neighbors=8, p=2, weights='uniform')`

In [0]: `plt.plot(n_range, clf_gs.cv_results_['mean_test_score'])`

Out[0]: `[<matplotlib.lines.Line2D at 0x7f13d6f41940>]`



3.6. Сравнение качества обучения моделей

```

In [0]: knn_best_model = KNeighborsClassifier(n_neighbors=24)
        knn_best_model.fit(X_train, Y_train)
        predicted_best = knn_best_model.predict(X_test)
        predicted_best

```

Out[0]: `array([1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0])`

```

In [0]: (accuracy_score(Y_test, predicted),
        accuracy_score(Y_test, predicted_best))


```

Out[0]: `(0.5737704918032787, 0.6557377049180327)`

```
In [0]: plot_confusion_matrix(Y_test, predicted,
                               classes=np.array(['E', 'E10+', 'M', 'T']),
                               title='Confusion matrix, without normalization')
```

Confusion matrix, without normalization


Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f13d6f4a320>



```
In [0]: plot_confusion_matrix(Y_test, predicted_best,
                               classes=np.array(['E', 'E10+', 'M', 'T']),
                               title='Confusion matrix, without normalization')
```

Confusion matrix, without normalization

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f13d6eeacf8>



```
In [0]: (precision_score(Y_test, predicted, average='weighted'),
        precision_score(Y_test, predicted_best, average='weighted'))
```

```
Out[0]: (0.5753212228622065, 0.6591074681238616)
```

```
In [0]: (recall_score(Y_test, predicted, average='weighted'),
        recall_score(Y_test, predicted_best, average='weighted'))
```

```
Out[0]: (0.5737704918032787, 0.6557377049180327)
```

```
In [0]: (f1_score(Y_test, predicted, average='weighted'),
        f1_score(Y_test, predicted_best, average='weighted'))
```

```
Out[0]: (0.5688953176899175, 0.65293502680339)
```

Таким образом, подбор гиперпараметра позволил улучшить результаты оценки модели посредством всех представленных метрик.

3.7. Кривые обучения и валидации

```
In [0]: from sklearn.model_selection import learning_curve, validation_curve
```

```
def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
```

```

return plt

def plot_validation_curve(estimator, title, X, y,
                          param_name, param_range, cv,
                          scoring="accuracy"):
    train_scores, test_scores = validation_curve(
        estimator, X, y, param_name=param_name, param_range=param_range,
        cv=cv, scoring=scoring, n_jobs=1)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

    plt.title(title)
    plt.xlabel(param_name)
    plt.ylabel("Score")
    plt.ylim(0.0, 1.1)
    lw = 2
    plt.plot(param_range, train_scores_mean, label="Training score",
             color="darkorange", lw=lw)
    plt.fill_between(param_range, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.2,
                     color="darkorange", lw=lw)
    plt.plot(param_range, test_scores_mean, label="Cross-validation score",
             color="navy", lw=lw)
    plt.fill_between(param_range, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.2,
                     color="navy", lw=lw)
    plt.legend(loc="best")
    return plt

```

```

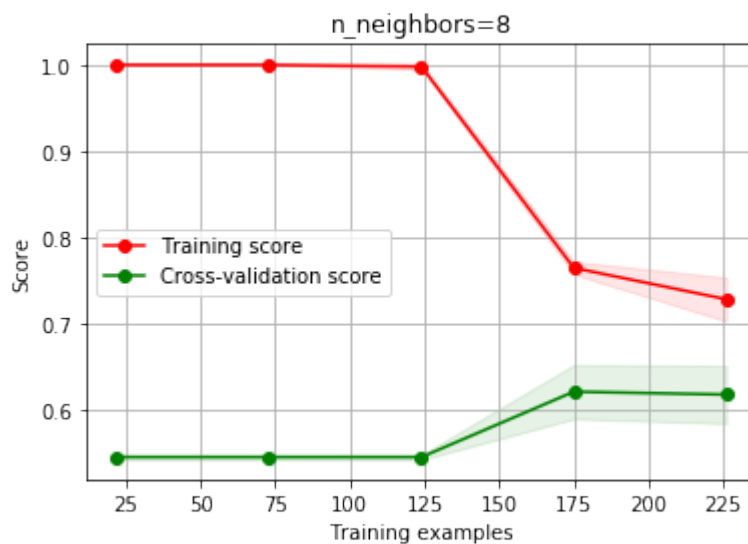
In [0]: plot_learning_curve(KNeighborsClassifier(n_neighbors=8),
                             'n_neighbors=8',
                             data_cleared,
                             target,
                             cv=4
                             )

```

```

Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py

```



```
In [0]: n_range2 = np.array(range(2,24,2))
```

```
plot_validation_curve(KNeighborsClassifier(), 'knn',
                      data_cleared,
                      target,
                      param_name='n_neighbors',
                      param_range=n_range2,
                      cv=4,
                      scoring="accuracy"
                      )
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py'
```

